Neural Networks

CS212 Topics in Computing 2 Module 2 Lecture 4

Spring 2018

Dr. Dmitri Roussinov dmitri.roussinov@strath.ac.uk

Lab 2

- No mass glitches
- Everybody survived

Big picture so far

- Learned: derivatives and gradients are crucial to train the models by optimizing their accuracy (reducing errors)
- Learned: how vector dot products can help to express the model in simpler and more efficient way
- Learned: how we can use Theano to efficiently obtains the gradients
- Just started: how we can use that to create an Artificial Neural Network so solve real life applications
 - E.g. handwritten digits recognition

Dot Products Review

Lab 3 Tips

- Tasks 10-11 are explained in Lecture 3 slides
- Task 12 expects you to read from a file
 - Template posted
 - Asks to use different variable sizes

Possible Error: Size Mismatch

```
L3-12 answer.py × L3-12 answer - size mismatch.py × basic.py ×
L3-12 answer.pv
                                 #import ...
External Libraries
                                   input vector = theano.tensor.fvector('input vector')
                                  torget values = theano.tensor.fvector('target values
                                  W = theano.shared(numpy.zeros((3,2)), 'W')
                                   \#W = \text{theano.shared(numpy.zeros((3,4)), 'W')}
                                  activations = theano.tensor.dot(W, input vector)
                         10
                                  predicted values = theano.tensor.nnet.sigmoid(activa
L3-12 answer - size mismatch
Debugger | ■ Console →
           "C:\Users\xepu8186\Appuata\Loca1\Cont1nuum\M1n1condaZ\11b\s1te-packages\tneano\go1\op.py
        r = p(n, [x[0] for x in i]
            C:\Users\xeb08186\AppData\Local\Continua.\Miniconda2\lib\site-packages\theano\tensor\basi
        z[0] = numpy.asarray(numpy.dot(x, y))
    ValueError: shapes (3,2) and (4,) not aligned: 2 (dim 1) = 4 (dim 0)
    Apply node that caused the error: dot(W, input vector)
    Topose index: 1
    Inputs types: [TensorType(float64, matrix), TensorType(float32, vector)]
    Inputs shapes: [(3L, 2L), (4L,)]
    Inputs strides: [(16L, 8L), (4L,)]
    Inputs values: ['not shown', array([ 0., 1., 0., 0.], dtype=float32)]
    Outputs clients: [[Elemwise (Scalar Sigmoid) [(0, 0)] (dot 0)]]

Dmitri Roussinov, CIS University of Strathclyde
```

Possible Error: Input/Output Mismatch

```
L3-12 answer.py × L3-12 answer - input mismatch.py ×
L3-12 answer.py
External Libraries
                       52
                                # training:
                                for epoch in xrange(1000): #run each training example 10 times #explair.
                       53
                                         for labely, vector in data train: labely: <type 'list'>: [0, 0,
                       54
                                              Accuracy, predicted values = train(vector, labely)
                       55
                                              if epoch % 100 == 0: #print progress on every 100-th epoch
                       56
                                               print epoch, Accuracy, predicted values, labely, vector
                       57
   L3-12 answer - inputs mismatch
Debugger 🔳 Console →" 🛌
   C:\Users\xeb08186\AppData\Local\Continuum\Miniconda2\python.exe "C:\Program Files (x86)\JetBrains\PyCharm 2016.3.3\helpers
   pydev debugger: process 9988 is connecting
   Connected to pydev debugger (build 163.15188.4)
   added to data: [0, 0, 1] [0, 1, 0, 0]
   added to data: [0, 0, 1] [1, 0, 0, 0]
   added to data: [0, 1, 0] [0, 0, 0, 0]
   added to data: [1, 0, 0] [0, 0, 1, 0]
   added to data: [1, 0, 0] [0, 0, 0, 1]
   Traceback (most recent call last):
     File "C:\Program Files (x86)\JetBrains\PyCharm 2016.3.3\helpers\pydev\pydevd.py", line 1596, in <module>
       globals = debugger r
                                                     _is module)
          rrogram Files (x86)\JetBrains\PyCharm 2016.3.3\herper 'dev\pydevd.py", line 974, in run
       pydev imports.execfile(file, globals, locals) # execute the script
     File "C:/H/NonRestorable/strath/212/Labs/L3-12 answer - input mismatch.py", ine 55, in <module>
       Accuracy, predicted values = train(vector, labely)
   ValueError: need more than 1 value to unpack
```

Lab 3-14: Last Part

- Using a Neural Network to implement handwritten digits recognition
- Separately, due to snow cancellations 2 weeks back
- Due before our test time
- Extra sessions this week (I will post times)
 - You can get help to wrap up
 - You can get it marked
- Right after the test you can get Lab 3 (all parts) marked but
 - We may not be able to give you enough help
 - We may be not be able to mark all so some will be marked base on your uploads
 - So do you best to show at extra sessions to get more feedback and help!

Lab 3-14: Last Part Competition Format

Let's try again a few demos

- vqa.cloudcv.org
- http://alteredqualia.com/xg/examples/morp hs_paintings.html

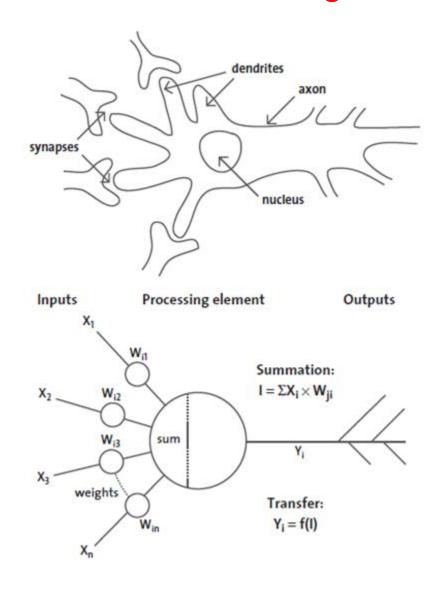
Plan for Today

- A few tips on Lab 3
- More about Neural Networks
 - Application to recognizing MNIST digits
- A few tips on last part Lab 3-14
- Possibly: questions about the sample test
- Wrapping up!

High Level Steps when Optimizing (Training) using Theano

- Define your prediction model
 - e.g. linear y = w1*x1+w2*x2
- Define your accuracy
 - e.g. negative sum of square errors
- Obtain the gradient
 - e.g. *G*=theano.tensor.grad(y, [w1,w2])
- Define updates for the training function
 - e.g. [(w1, w1+.1*G[0]), (w2, w2+.1*G[1])]
- Call the training function in the loop
- Done!

Artificial Neural Network: Biological Motivation



Dmitri Roussinov, CIS University of Strathclyde

















Things To Note

- Input variables are converted into output variables
- One node takes several inputs and produces one output
- Linear conversion with weights can be used in the simplest case
- Outputs can be used as inputs for other nodes
- Thus, multiple layers can be implemented
- Function f provides non-linearity
- Most often it is sigmoid









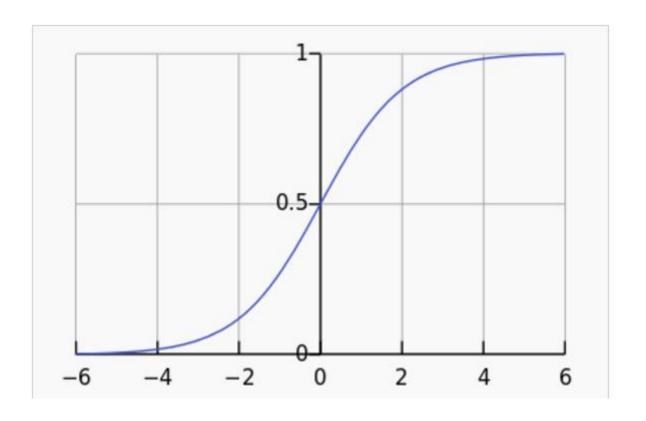








Non linearity: sigmoid function



$$S(t) = \frac{1}{1+e^{-t}}.$$



















Things To Note

- First, the value is calculated by adding all inputs with weights
- Then, the sigmoid function is applied to the value:
 - For large positive values sigmoid approaches 1
 - For large negative values sigmoid approaches 0
 - For intermediate values, the transition from 0 to 1 is a smooth mathematical function
 - That allows a number of simplifications when obtaining its gradient











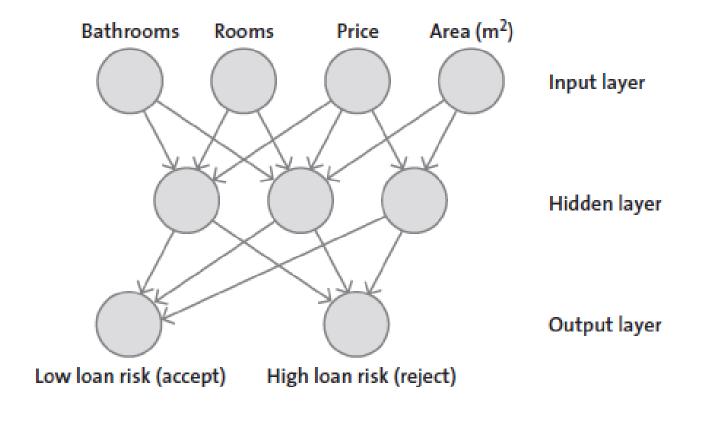






ANN: Architecture

Figure 5.13 The architecture of a artificial neural network for a hypothetical predictor of domestic property loan risk



Terminology

- Each node is called a neuron
- Neurons are organized into layers
- Normally, neurons are connected to other neurons only on adjacent layers
- Input layer is the one that represents the features (AKA independent variables)
 - E.g. #bathrooms, price, area
- Output layer is the one that represents the predictions (AKA label, AKA dependent variable)
 - E.g. accept/reject loan decision
- Hidden is a layer that is neither input nor output















Backpropagation

- Gradient Ascent/Descent is a common method of training artificial neural networks so as to minimize the prediction errors
- When applied to Neural Networks, it is called Backpropagation
- We already know how use it in general case and define the model in **Theano**
- That saves out a lot of hassle

















MNIST task



What are the challenges?

















Overall Template Posted: 2 Screens of code

- Templates and Test/Exam:
 - The details from the next two slides are not directly tested
 - The program code from the template posted for lab 3 is not directly tested
 - All other slides are testable unless excluded

Screen 1: Dummy Theano model and reading data

```
import ...
input vector = theano.tensor.fvector('input vector') #theano variable representing image
target values = theano.tensor.fvector('target values') #theano variable representing the I
#This is just a template: it does not learn anything, and always returns the digit "0":
params = []
updates = []
theano.config.on unused input='ignore'
predicted class = theano.shared(0, 'predicted class')
Accuracy = theano.shared(0, 'Accuracy')
#Change this to something meaningful and it will work!
# defining Theano functions for training and testing the model:
train = theano.function([input vector, target values], [Accuracy, predicted class], update
test = theano.function([input vector, target values], [Accuracy, predicted class], allow i
    #'allow input downcast=True' is needed to avoid any issues converting between 64 and 3
def read dataset(path): #The function that reads training or testing data and returns it a
    number of images = len(open(path).readlines()) / 28 #28 lines per each image
    f = open(path)
    dataset = [] #starts with an empty container
    for i in range(number of images):
     data vector = [] #we start with empty data vector, and then we read the data for each
     for 1 in range (28): #each image takes 28 lines
```

Screen 2: training/testing

```
#reading the data:
data test = read dataset("test.txt")
data train = read dataset("train.txt")
# training
for epoch in xrange(10):
        cost sum = 0.0
        correct = 0
        for labely, vector in data train:
            Accuracy, predicted class = train(vector, labely)
            cost sum += Accuracy
            if (labelv[predicted class] == 1):
                correct += 1
        print "Epoch: " + str(epoch) + ", Accuracy: " + str(cost sum) + ", %correct: " + str(flo
# testing:
cost sum = 0.0
correct = 0
for labely, vector in data test:
        Accuracy, predicted class = test(vector, labely)
        cost sum += Accuracy
        if (labelv[predicted class] == 1):
                    correct += 1
print "\t%correct on the test set: " + str(float(correct) / len(data test)*100)
```

Things to Note

- Outline similar to the one for Task 12 (Sentiment from data in a file)
 - Added Theano function 'test' to test the trained accuracy
 - It is almost same as 'train'
 - It calculate the predictions but does not obtain gradients and does not change W
- 'read_dataset' function similar to Task 12, but hints:
 - Input format a bit different (explained below)
 - Sizes are now different

Output: Nothing Useful Yet

```
L3 template.py × bydevd.py ×
                          проси.
                                      · DCI (CDOCII)
        # testing:
        cost sum = 0.0
        correct = 0
51
        for labely, vector in data test:
53
                  Accuracy, predicted class = test(
54
                  cost sum += Accuracy
                  if (labelv[predicted class] == 1)
  L3 template
ebugger 🔳 Console → 🖢
                    王义坐义之为
   C:\Users\xeb08186\AppData\Local\Continuum\Miniconda2\pythc
   pydev debugger: process 12048 is connecting
   Connected to pydev debugger (build 163.15188.4)
   Epoch: 0, Accuracy: 0.0, %correct: 9.94444444444
   Epoch: 1, Accuracy: 0.0, %correct: 9.94444444444
   Epoch: 2, Accuracy: 0.0, %correct: 9.94444444444
   Epoch: 3, Accuracy: 0.0, %correct: 9.94444444444
  Epoch: 4, Accuracy: 0.0, %correct: 9.94444444444
   Epoch: 5, Accuracy: 0.0, %correct: 9.94444444444
   Epoch: 6, Accuracy: 0.0, %correct: 9.94444444444
   Epoch: 7, Accuracy: 0.0, %correct: 9.94444444444
   Epoch: 8, Accuracy: 0.0, %correct: 9.94444444444
   Epoch: 9, Accuracy: 0.0, %correct: 9.94444444444
       %correct on the test set: 8.5
   Process finished with exit code 0
```

Dr

More to Notice:

- One epoch finishes when each training example is fed once
- We print accuracy (%correct) after each epoch
- After 10 epochs, the overall accuracy (%correct) is computed over a dataset that the model has not seen called a test-set
- This is a standard procedure to simulate real-life scenario
- Currently it is 10% since our model always tells it is the digit "0"

Data Format

Things To Note:

- 28 lines per image
- Label (digit) is the very first number on each line
- 9000 images to train
- 1000 images to test

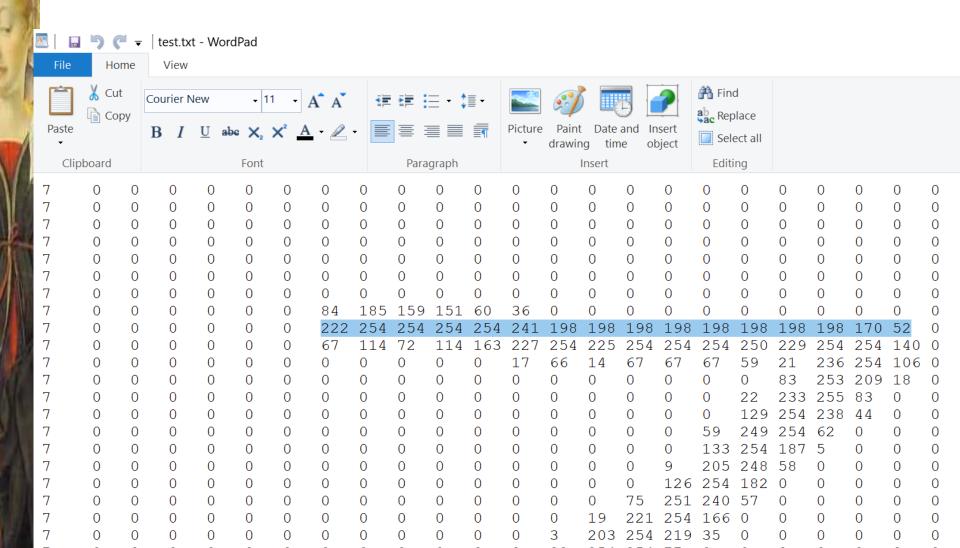
The code to create one-hot vectors in the template

```
def read dataset (path): #The function that reads training or testing data
   number of images = len(open(path).readlines()) / 28 #28 lines per each
   f = open(path)
   dataset = [] #starts with an empty container
    for i in range (number of images):
     data vector = [] #we start with empty data vector, and then we read :
     for 1 in range(28): #each image takes 28 lines
            line = f.readline()
            line parts = line.split() #split the line into all the number:
            assert(len(line parts) == 29) #should be total of 29: the labe
            label = int(line parts[0]) #very first number in the file is
            assert (0 <= label <= 9) #only digits 0-9 are allowed as labe.
            #nov; we will create "one-hot vector":
            label vector = [0,0,0,0,0,0,0,0,0] #a vector of 10 2 roes
            label vector[label] = 1 #except 1 for the label
            data vertor += [float(line parts[i])/255 for i in xrange(1, i
            #we divide by 255 so the pixel brightness is represented by a
```

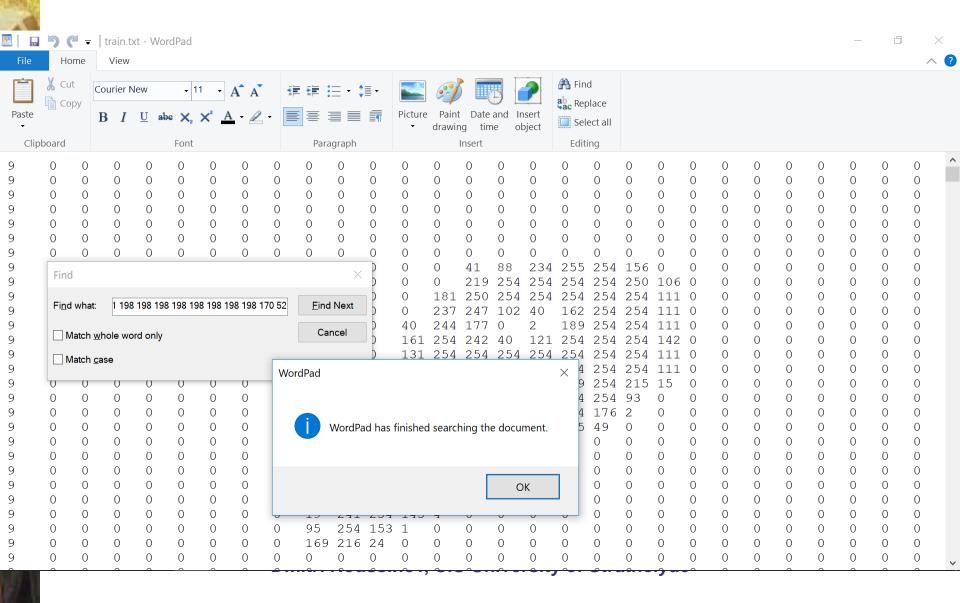
dataset.append((label_vector, data_vector))
return dataset

How to verify train/test don't overlap (Task 14-a)

Select some "non trivial" line in test



Search for it in train

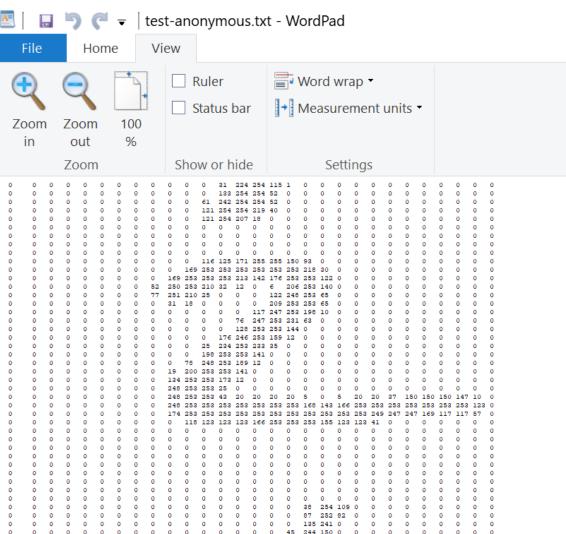


What to check

- If there is not such line, we are fine
- Can check a few more randomly selected sequences
- If we find it, verify that the entire image is still not the same

Let's do altogether

Testing without giving the labels



Still Works!

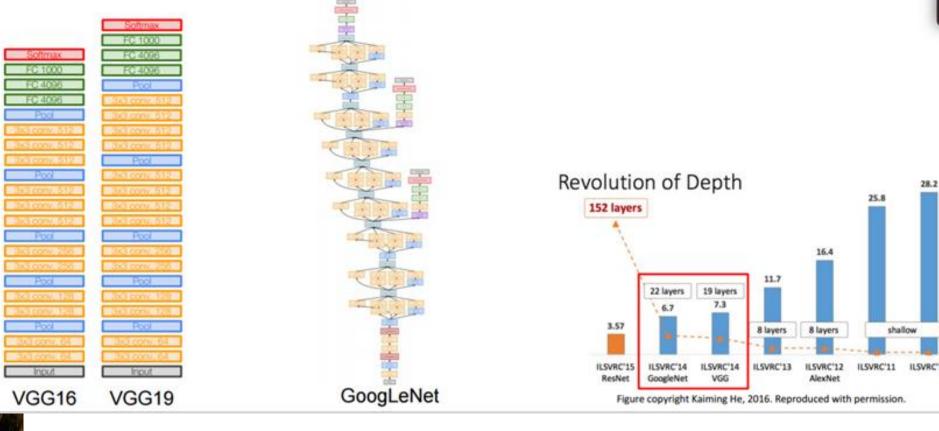
```
H > MonRestorable > m strath > m 212 > m Labs > 6 L3 template-MNIST-anon.py >
                      L3 template-MNIST.py × 2 L3 template-MNIST-to-solution.p
12 answer.py
ernal Libraries
                      63
                               # testing:
                               cost sum = 0.0
                      64
                               correct = 0
                      65
                               for labely, vector in data test:
                      66
                                    Accuracy, predicted class = 1
                      67
                                    cost sum += Accuracy
                      68
                     69
                                    if (labelv[predicted class] :
                                          correct += 1
                     70
 L3-12 answer - inputs mismatch
                            L3 template-MNIST-anon
gger | ■ Console → ■ 🛌
 C:\Users\xeb08186\AppData\Local\Continuum\Miniconda2\python.exe "C:\Proc
pydev debugger: process 2204 is connecting
 Connected to pydev debugger (build 163.15188.4)
 Epoch: 0, Accuracy: -2254.18465691, %correct: 86.1222222222
 Epoch: 1, Accuracy: -1742.21647648, %correct: 89.944444444
 Epoch: 2, Accuracy: -1599.65213381, %correct: 90.977777778
```

Let's try to create more layers

And do a question from the sample

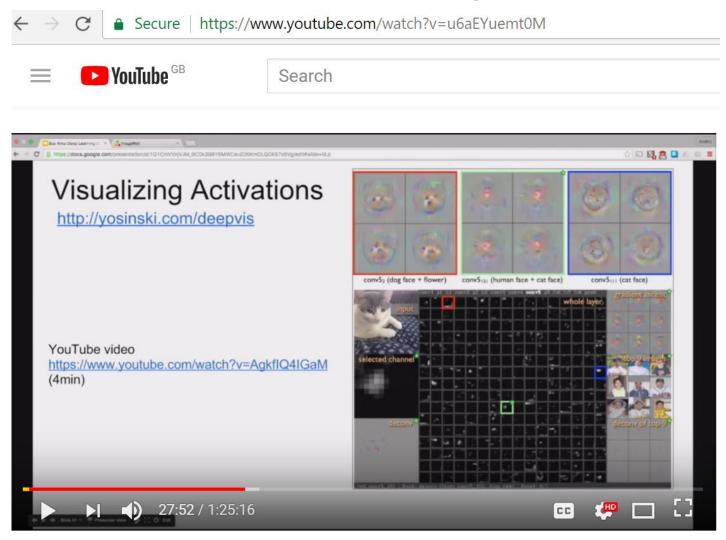
Modern Computer Vision Architecture

Last Time: CNN Architectures



For illustration only: Not to be tested

How this works in larger networks



Deep Learning for Computer Vision (Andrei Karpathy, OpenAI)

(for illustration only)